

Center for Bioinformatics &  
Computational Biology

BIOINFORMATICS  
CORE FACILITY

# High Performance Computation (HPC) Basics

FAIR Data Practices for Omics Analysis Workshop

University of Delaware

April 19 (Day 2)

# Starting with BASH

- BASH is used by issuing commands at the prompt
- The prompt is always at a location in the file system
  - When you login you should be in your "home folder": /home/username
  - Your home directory is where you can store your own files (and install software)
- You can run a command by typing its name at the prompt and hitting [enter/return]
- Some commands usually have options, called flags or parameters
  - Flags can stand by themselves.  
"ls -l ": ls is a command and -l is a flag
  - Or they can require arguments  
"tar -c -f file.tar input-dir":
    - -c and -f are flags
    - file.tar is the argument for -f.
    - -c does not need an argument.
    - input-dir is a standalone argument. Some commands expect these and they almost always indicated what the input file (or directory) for a command is. They are usually the last thing listed in a command.

# BASH Essentials

(90% of what you need to know in 1 slide)

- Command prompt is used to navigate the file system and is always set in a directory [*pwd*]
- Directory can be created [*mkdir*], changed [*cd*], or removed [*rmdir*]. Or its contents can be listed [*ls*]
- Files and directories have read, write, and execute permissions that can be set at the owner, group, and system levels [*chmod*, *chown*, *chgrp*]
- Files can be found [*find*], displayed [*cat*, *head*, *tail*, *less*], searched [*grep*], or edited/created [*nano*, *vi*, *emacs*]
- Scripts and binary applications with the execute permission can be run by typing their name
- Help about any command can be displayed [*man*] to learn about the syntax and options (parameters & flags)
- Standard output often goes to the screen by default, but can be saved (redirected) into a file using the > (replace) or >> (append) operators, or sent as input to another command using the | (pipe) operator
- Remember to use [*man*] to get help about almost any system command
- TIP: Spend some time learning to use the very powerful commands *grep*, *sed*, and *awk*

# What is HPC?

# HPC = High Performance Computing

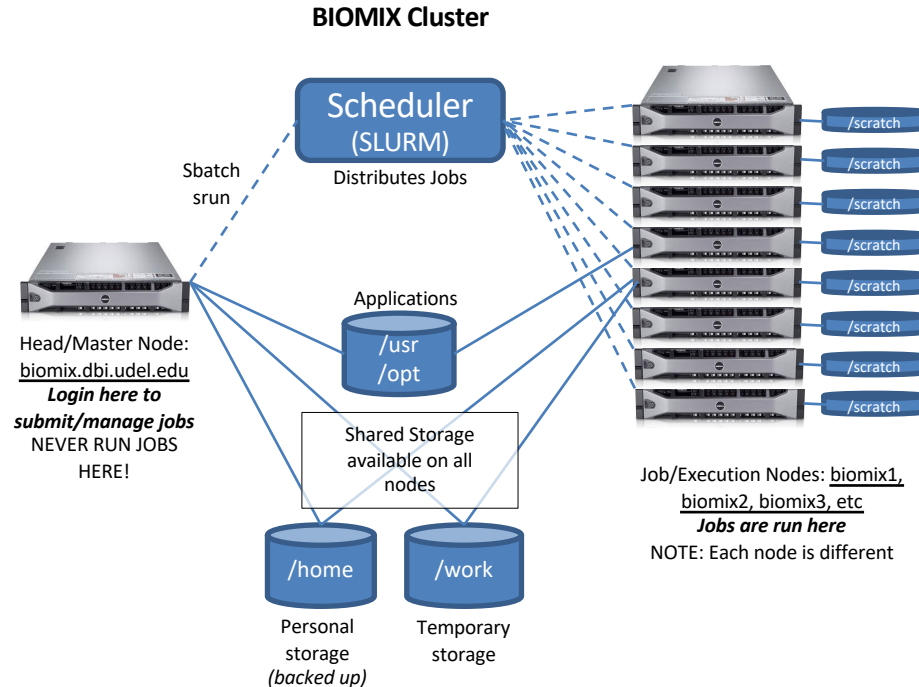
# BIOMIX Cluster

- Co-hosted by CBCB and DBI
- Heterogeneous cluster with >2750 cpu cores
- Each node has 12-96 CPU cores and 128-2048GB of RAM and some amount of /scratch storage
- Each user has a home directory (storage quota)
- Username and password (change password with this: *passwd*)
- <https://bioit.dbi.udel.edu> (BIOMIX manual and software)
- <https://bioit.dbi.udel.edu/BIOMIX/SLURM-intro.html> (slurm info)

# Biomix Nodes



# High Performance Computing (HPC) Cluster





# Interactive Submission

- Never run an application on the head node
- Use interactive submission when you want to experiment or manually run something
- `srun -c 1 --mem=16000 --pty bash`
  - `-c 1` = reserve 1 cpu core
  - `--mem` = reserve 16000MB of memory (16GB)
  - `--pty` = run a pseudo-terminal
  - `bash` = run bash on the remote node
- IMPORTANT – When Finished Type: *exit*

# Batch Submission

- Allows one or more commands to be submitted for automatic execution
- This is the typical way things are run on the cluster
- *sbatch script.slurm*
- By default STDOUT and STDERR goes to slurm-[jobid].out (we will change this a bit in our jobs)

# Batch Submission

## (Example slurm submission script)

```
#!/bin/bash
#SBATCH --job-name=blast
#SBATCH --mem=64000
#SBATCH -c 4 ← Number of cpu cores

#insert software requirements here

cd ~/exercises/blast-exercise ← Tells the command what directory to run in...

echo "Start:"
date
hostname

# PUT THE COMMAND(S) YOU WANT TO RUN NEXT
blastp -query /home/spolson/exercises/blast-exercise/sequence.fasta -db /usr/local/blast_db/nr
-out /home/spolson/exercises/blast-exercise/blast-result.txt -num_threads 4 -outfmt 6
                                     ↑
                                   Number of cpu cores

echo "Complete:"
date
```

NOTE: A slurm software update eliminated the "#SBATCH --workdir" option that appears in the lecture video. This is replaced by the "cd" command below.

# Monitoring Cluster, Queue, Jobs

- List nodes and their capabilities:
  - *scontrol show nodes*
  - *cat /etc/dbi/NodesDB*
- List current jobs running:
  - *queue*
- Details on a job 2992
  - *scontrol show job 2992*